

Bayesian neural networks for sparse coding

Danil Kuzin ¹ Olga Isupova ² Lyudmila Mihaylova ¹

¹Department of Automatic Control and Systems Engineering, University of Sheffield, UK

²Machine Learning Research Group, University of Oxford, UK

ICASSP 2019






Problem

Sparse regression

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon},$$

$$\mathbf{y} \in \mathbb{R}^K, \boldsymbol{\beta} \in \mathbb{R}^D, K < D$$

How to find sparse $\boldsymbol{\beta}$ given \mathbf{y}, \mathbf{X} ?







Sparse frequentist regression

Add sparsity-inducing penalty


$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \left[\|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \|\beta\|_p^p \right]$$




l_2 -penalty



$l_{1/2}$ -penalty



l_1 -penalty




$l_{1/8}$ -penalty



Sparse Bayesian regression

Add sparsity-inducing prior

Strong sparsity




$$\beta_d \sim (1 - z_d)\mathcal{N}(0, \sigma^2) + z_d\delta_0$$

$$z_d \sim \text{Ber}(\omega)$$

- probability of exact zero
- discrete variables

Weak sparsity



$$\beta_d \sim \mathcal{N}(0, \sigma_d^2)$$

$$\sigma_d^2 \sim \text{IG}(a)$$

- continuous at zero
- continuous variables



Problem


Build a model that estimates β from observations \mathbf{y} such that $\mathbf{y} = \mathbf{X}\beta + \varepsilon$, and elements of β contain zeros.



Iterative Shrinkage-Thresholding Algorithm (ISTA)

Soft thresholding

$$h_{\lambda}(\mathbf{b}) = \text{sgn}(\mathbf{b}) \max(|\mathbf{b}| - \lambda, 0),$$



Prediction

Require: observation \mathbf{y}

- 1: *Define.* weights $\mathbf{W} = \mathbf{X}^T / E$, E — the largest eigenvalue of $\mathbf{X}^T \mathbf{X}$ and $\mathbf{S} = \mathbf{I}_{D \times D} - \mathbf{W} \mathbf{X}$.
- 2: *Initialisation.* Dense layer $\mathbf{b} \leftarrow \mathbf{W} \mathbf{y}$
- 3: *Initialisation.* Soft-thresholding nonlinearity $\hat{\beta}_0 \leftarrow h_{\lambda}(\mathbf{b})$
- 4: **repeat**
- 5: Dense layer $\mathbf{c}_l \leftarrow \mathbf{b} + \mathbf{S} \hat{\beta}_{l-1}$
- 6: Soft-thresholding nonlinearity $\hat{\beta}_l \leftarrow h_{\lambda}(\mathbf{c}_l)$
- 7: **until** converged
- 8: **return** $\hat{\beta} \leftarrow \hat{\beta}_L$

Daubechies I., Defrise M., De Mol C. An iterative thresholding algorithm for linear inverse problems with a sparsity constraint. Communications on pure and applied mathematics, 2004.




Problem

Train a model on N data examples \mathbf{B}, \mathbf{Y} that estimates β from observations \mathbf{y} such that $\mathbf{y} = \mathbf{X}\beta + \epsilon$, and elements of β contain zeros.



Learned ISTA (LISTA)

Learn weights \mathbf{W}, \mathbf{S} based on training data



Forward pass

Require: observation \mathbf{y} , current weights \mathbf{W}, \mathbf{S} , current λ , number of layers L

- 1: *Initialisation.* Dense layer $\mathbf{b} \leftarrow \mathbf{W}\mathbf{y}$
- 2: *Initialisation.* Soft-thresholding nonlinearity $\hat{\beta}_0 \leftarrow h_\lambda(\mathbf{b})$
- 3: **for** $l = 1$ **to** L **do**
- 4: Dense layer $\mathbf{c}_l \leftarrow \mathbf{b} + \mathbf{S}\hat{\beta}_{l-1}$
- 5: Soft-thresholding nonlinearity $\hat{\beta}_l \leftarrow h_\lambda(\mathbf{c}_l)$
- 6: **end for**
- 7: **return** $\hat{\beta} \leftarrow \hat{\beta}_L$

Backward pass

- 1: **update** $\mathbf{W}, \mathbf{S}, \lambda$ based on derivatives of the mean squared loss

Gregor, K., LeCun, Y. Learning fast approximations of sparse coding. ICML 2010



BayesLISTA

Prior

$$p(\mathbf{W}) = \prod_{d=1}^D \prod_{k=1}^K \mathcal{N}(w_{ij}; 0, \eta^{-1}), \quad p(\mathbf{S}) = \prod_{d'=1}^D \prod_{d''=1}^D \mathcal{N}(s_{d'd''}; 0, \eta^{-1}),$$

Forward pass

Propagate the distribution for $\hat{\beta}$ through layers, add Gaussian noise

$$p(\beta|\mathbf{y}, \mathbf{W}, \mathbf{S}, \gamma, \lambda) = \prod_{d=1}^D \mathcal{N}\left(\beta_d; [\hat{\beta}(\mathbf{y}; \mathbf{S}, \mathbf{W}, \lambda)]_d, \gamma^{-1}\right)$$

Posterior

$$p(\mathbf{W}, \mathbf{S}, \gamma, \eta | \mathbf{B}, \mathbf{Y}, \lambda) = \frac{p(\mathbf{B}|\mathbf{Y}, \mathbf{W}, \mathbf{S}, \gamma, \lambda) p(\mathbf{W}|\eta) p(\mathbf{S}|\eta) p(\eta) p(\gamma)}{p(\mathbf{B}|\mathbf{Y}, \lambda)}$$

Backward pass


Update weights with probabilistic backpropagation




Uncertainty propagation: Initialisation

At every step the output of soft-thresholding can be closely approximated with the spike and slab distribution


1. $\mathbf{b} = \mathbf{W}\mathbf{y}$ is
Gaussian-distributed



2. $\hat{\beta}_0 = h_\lambda(\mathbf{b})$ is
approximated with the spike
and slab distribution




3. $\mathbf{e}_l = \mathbf{S}\hat{\beta}_{l-1}$ is approximated
with the Gaussian
distribution






Uncertainty propagation: Iterations

4. $\mathbf{c}_I = \mathbf{b} + \mathbf{e}_I$ is
Gaussian-distributed



5. $\hat{\beta}_I = h_\lambda(\mathbf{c}_I)$ is
approximated with the spike
and slab distribution



Advantages

All latent variables are modelled with parametrised distributions

We can apply approximate Bayesian inference methods



Probabilistic backpropagation

Posterior

$$p(\mathbf{W}, \mathbf{S}, \gamma, \eta | \mathbf{B}, \mathbf{Y}, \lambda) = \frac{p(\mathbf{B} | \mathbf{Y}, \mathbf{W}, \mathbf{S}, \gamma, \lambda) p(\mathbf{W} | \eta) p(\mathbf{S} | \eta) p(\eta) p(\gamma)}{p(\mathbf{B} | \mathbf{Y}, \lambda)}$$

Approximate posterior

$$q(\mathbf{W}, \mathbf{S}, \gamma, \eta) = \prod_{d=1}^D \prod_{k=1}^K \mathcal{N}(w_{dk}; m_{dk}^w, v_{dk}^w) \prod_{d'=1}^D \prod_{d''=1}^D \mathcal{N}(s_{d' d''}; m_{d' d''}^s, v_{d' d''}^s) \\ \times \text{Gam}(\gamma; a^\gamma, b^\gamma) \text{Gam}(\eta; a^\eta, b^\eta)$$

Assumed density filtering (ADF)

Iteratively add factors from the posterior p one-by-one in the factorised approximating distribution q and update the approximation q .

Hernandez-Lobato, J. M., Adams, R. Probabilistic backpropagation for scalable learning of Bayesian neural networks, ICML 2015



Probabilistic backpropagation

Updates for Gaussian parameters

$$q(a) = Z^{-1} f(a) \mathcal{N}(a; m, v)$$

For each factor $a \in \{w_{dk}, s_{d' d''}\}$ new parameters are:

$$m \leftarrow m + v \frac{\partial \log Z}{\partial m}, \quad v \leftarrow v - v^2 \left[\left(\frac{\partial \log Z}{\partial m} \right)^2 - 2 \frac{\partial \log Z}{\partial v} \right]$$

$\frac{\partial \log Z}{\partial m}$, $\frac{\partial \log Z}{\partial v}$ are required

Updates for Gamma parameters

For each factor $a \in \{\eta, \gamma\}$ new parameters are:

$$\alpha \leftarrow [Z_0 Z_2 Z_1^{-2} (\alpha + 1) / \alpha - 1]^{-1}, \quad \beta \leftarrow [Z_2 Z_1^{-1} (\alpha + 1) / \beta - Z_1 Z_0^{-1} \alpha / \beta]^{-1}$$

$Z_0 = Z(\alpha, \beta)$, $Z_1 = Z(\alpha + 1, \beta)$, $Z_2 = Z(\alpha + 2, \beta)$ are required



Probabilistic backpropagation

The normalising constant when incorporating the likelihood factor:

$$Z = \int \prod_{d=1}^D \mathcal{N}(\beta_d; [\hat{\beta}(\mathbf{y}; \mathbf{s}, \mathbf{w}, \lambda)]_d, \gamma^{-1}) q(\mathbf{w}, \mathbf{s}, \gamma, \eta) d\mathbf{w} d\mathbf{s} d\gamma d\eta$$


Assuming the spike and slab distribution for $\hat{\beta}$:

$$Z \approx \prod_{d=1}^D \left[\omega_d^{\hat{\beta}} \mathcal{T}(\beta_d; 0, \beta^\gamma / \alpha^\gamma, 2\alpha^\gamma) + (1 - \omega_d^{\hat{\beta}}) \mathcal{N}(\beta_d; m_d^{\hat{\beta}}, \beta^\gamma / (\alpha^\gamma - 1) + v_d^{\hat{\beta}}) \right],$$


where $\{\omega_d^{\hat{\beta}}, m_d^{\hat{\beta}}, v_d^{\hat{\beta}}\}$ are the parameters of the spike and slab distribution for $[\hat{\beta}]_d$.



Synthetic Experiments



Different depth performance




Different observation size performance

$$N_{\text{train}} = 1000, D = 100, p(\beta_d = 0) = 0.8, \eta^{-1} = 0.25, \lambda = 0.1$$


MNIST Experiments

\mathbf{X} - random i.i.d, $D = 784$, $K = 100$


Posterior parameters for an image of digit 7



β posterior mean





β posterior std



β posterior spike indicator

Samples from the posterior for an image of digit 7






Active Learning

Idea

Use the estimated uncertainty to choose next training data with largest variance

$$N_{\text{train}} = 50, N_{\text{pool}} = 500, N_{\text{addition}} = 1$$





Contributions and future work

Key contributions

- Bayesian version of LISTA
- Uncertainty propagation to make inference feasible
- Uncertainty quantification, for example, for active learning

Future work

- The shrinkage parameter λ as a parameter of the model
- Stochastic inference

Code will be available at danilkuzin.github.io